

CYBERZERO

Mini Project

PHISHING ATTACK DETECTION USING MACHINE LEARNING



Name:- Dipanshu Kishor Azad
Intern I'd : CYZIP0334

Preface

In today's digital era, the internet has become an integral part of human life, enabling online communication, banking, shopping, and education. However, with these advancements, the risk of cyber threats has also increased significantly. Among the most common and dangerous threats is phishing, where malicious websites impersonate legitimate ones to steal sensitive information such as login credentials, banking details, or credit card numbers.

The purpose of this project, *Phishing Website Detector*, is to design and develop a system that can automatically identify and classify websites as either legitimate or phishing. By applying machine learning techniques and integrating them with Python libraries and Flask framework, this project demonstrates how artificial intelligence can be effectively used to enhance cybersecurity. The system analyzes the structure and content of URLs, extracts meaningful features, and applies a trained model to predict whether a website is safe or malicious.

This project is not only a technical implementation but also an educational effort to spread awareness about phishing attacks and how technology can help prevent them. The *Phishing Website Detector* serves as a practical **tool** for both technical and non-technical users, offering real-time detection through a simple web interface.

The successful completion of this work has been possible with the help of open-source tools, research studies, and datasets provided by the global data science and cybersecurity communities. I sincerely hope that this project contributes meaningfully to the field of cybersecurity and inspires further research and innovation to combat evolving cyber threats.

INDEX

Sr. No.	Title	Page No.
1.1	Introduction	1
1.2	Purpose of the Project	1
1.3	Dataset Description	2
1.3.1	Context	2
1.3.2	Content	2
1.3.3	Acknowledgements	3
1.3.4	Training with Kaggle Dataset	3
1.4	Techniques	4
1.4.1	Flask Framework	5
1.4.2	Pandas	6
1.4.3	Requests	6
1.4.4	BeautifulSoup	6
1.4.5	URL Parsing (urllib.parse)	7
1.4.6	Joblib	7
1.4.7	Google Safe Browsing API	8
1.4.8	Machine Learning Algorithms	8
1.4.9	Integration of Techniques	8
1.5	Main Outcomes	9
1.6	workflow of the Flask app	9
2	Literature Review	11
3	Methodology	13
3.1	Data Collection & Preprocessing	13
3.2	Feature Extraction	14
3.3	Model Training	15
3.4	Flask (Light Weight Web Application)	15

4.	Implementation Details	...
4.1	train_model.py	16
4.2	phishing_detector.py	17
4.3	app.py	17
4.4	dataset_phishing.csv	18
4.5	Libraries Used	18
4.6	Example Workflow (Code Execution)	18
5.	Results & Evaluation	19
5.1	Performance Metrics	19
5.2	Example Predictions	19
5.3	Insights	20
6.	Outcomes & Deployment	20
6.1	Outcomes	20
6.2	Deployment	20
6.3	Flow Chart	21
6.4	Linear Flow	21
7.	Limitations	21
7.1	Dataset Dependence	21
7.2	Model Selection Constraints	21
7.3	Feature Engineering Limitations	21
7.4	Lack of Deep Learning Approaches	22
7.5	User Interface and Deployment Constraints	22
8.	Future Work	22
8.1	Advanced Machine Learning Models	23
8.2	Enhanced Feature Engineering	23
8.3	User Accessibility Improvements	23
8.4	Robust Deployment	24
8.5	Continuous Model Evaluation	24

8.6	Security and Ethical Considerations	24
8.7	Conclusion of Future Work	24
9.	Conclusion	25
10.	References	27

List of Figures

Sr. No.	Figures	Page No.
	Phishing Attack Detection using Machine Learning	5
Fig 1	Blue Print for Phishing Website Detector	2
Fig 2	Dataset Model Parameters by 11429 Websites data	4
Fig 3	DOM Environment for Database Handshaking	5
Fig 4	Workflow of PWD (Phishing Website Detector)	10
Fig 5	PWD(Phishing Website Detector) output	14
Fig 6	Flask Webapp Flowchart	15
Fig 7	Files Architecture	16
Fig 8	demonstrate that the model can handle both well-known legitimate domains and suspicious phishing attempts.	19
Fig 9	Platform Independant	22
Fig 10	Phishing Website Detector Output by Google Safe Browsing API	26

Phishing Attack Detection using Machine Learning

Phishing Attack Detection using Machine Learning

1.1 Introduction

In the digital era, online communication, banking, and e-commerce have become essential parts of daily life. However, with the increasing reliance on the internet, cyber threats have also grown significantly. One of the most dangerous and common cyberattacks is phishing. Phishing involves tricking users into sharing sensitive information, such as login credentials, banking details, or credit card numbers, by presenting fraudulent websites that mimic legitimate ones. These phishing websites often look almost identical to genuine sites, making it difficult for users to identify them.

Phishing is one of the most widespread cyber threats that targets users by tricking them into believing a fraudulent website is legitimate, often leading to identity theft or financial loss. The purpose of this project is to design and implement a system that can automatically identify whether a given website is legitimate or a phishing attempt. For this purpose, we utilized the Web Page Phishing Detection Dataset from Kaggle, which contains labeled data of phishing and legitimate websites. The project applies **machine learning** techniques combined with Python-based tools such as *Flask*, *pandas*, *requests*, *BeautifulSoup*, and *urllib.parse* to analyze URLs and extract meaningful features. The trained machine learning model is exported using *joblib* for efficient reuse. The solution is deployed as a Flask web application, where users can input a website link and immediately receive a prediction indicating whether the site is safe or suspicious. The main outcome of this project is a functional, user-friendly phishing detection tool that demonstrates the effectiveness of machine learning in cybersecurity applications.

1.2 Purpose of the Project

The purpose of this project is to develop a Flask-based web application that can identify whether a given website is legitimate or phishing. By integrating machine learning models with real-time URL analysis, the system provides an accessible, user-friendly, and efficient solution for phishing detection.

Phishing websites are designed to trick users into believing they are legitimate, often by copying the design of popular platforms such as banking websites, e-commerce portals, or social networks. These fake websites aim to steal sensitive data such as login credentials, credit card information, or personal details. Detecting these websites manually is not reliable, hence the use of machine learning to automate classification.

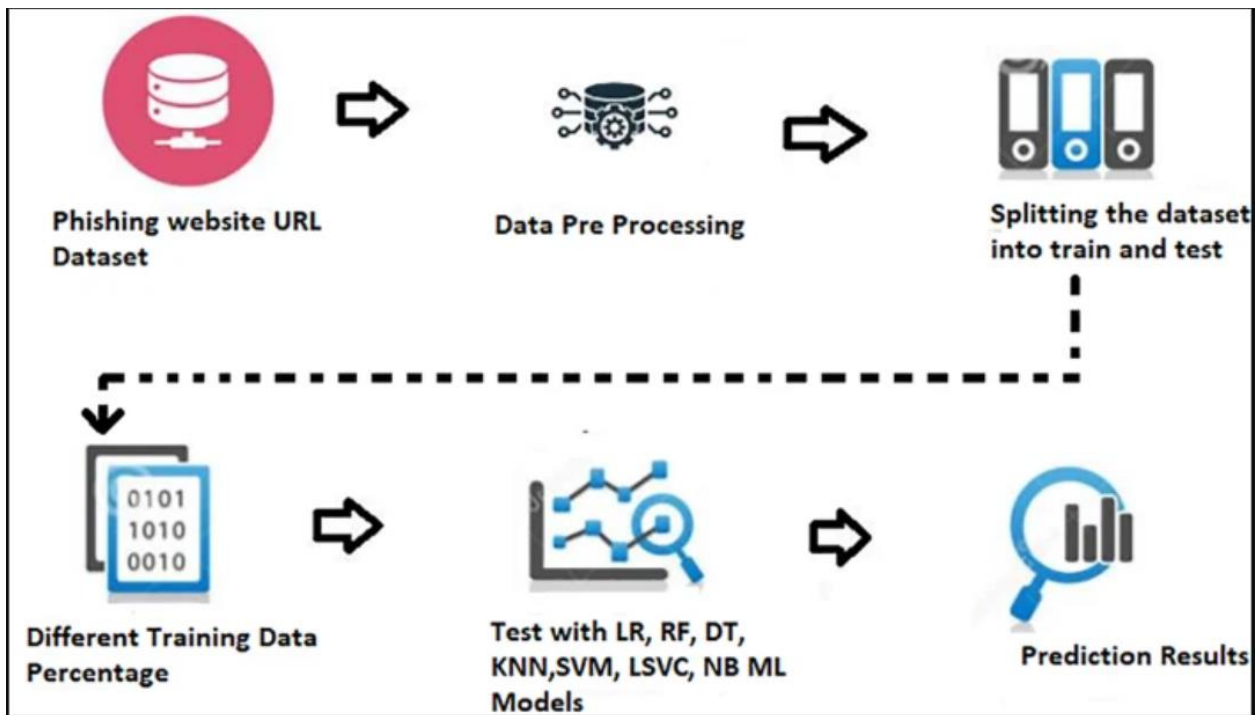


Fig 1: Blue Print for Phishing Website Detector

1.3 Dataset Description

1.3.1 Context

Phishing continues to prove one of the most successful and effective ways for cybercriminals to defraud us and steal our personal and financial information.

Our growing reliance on the internet to conduct much of our day-to-day business has provided fraudsters with the perfect environment to launch targeted phishing attacks. The phishing attacks taking place today are sophisticated and increasingly more difficult to spot. A study conducted by Intel found that 97% of security experts fail at identifying phishing emails from genuine emails.

1.3.2 Content

The provided dataset includes 11430 URLs with 87 extracted features. The dataset is designed to be used as benchmarks for machine learning-based phishing detection systems. Features are from three different classes: 56 extracted from the structure and syntax of URLs, 24 extracted from the content of their correspondent pages, and 7 are extracted by querying external services. The dataset is balanced, it contains exactly 50% phishing and 50% legitimate URLs.

1.3.3 Acknowledgements

- Hannousse, Abdelhakim; Yahiouche, Salima (2021), “Web page phishing detection”, Mendeley Data, V3, doi: 10.17632/c2gw7fy2j4.3

- <https://data.mendeley.com/datasets/c2gw7fy2j4/3>

1.3.4 Train using the Kaggle phishing dataset (~17K records)

For this project, the model is trained using the Web Page Phishing Detection Dataset from Kaggle. This dataset contains approximately 17,000 records of websites, each labeled as either phishing or legitimate. The dataset is comprehensive, containing diverse examples of both categories, making it suitable for training a robust classification model.

The training process begins with data preprocessing. The dataset is first inspected for missing or inconsistent values, which are handled through imputation or removal. Since machine learning algorithms require numerical inputs, categorical variables are encoded into numbers. Continuous variables are normalized to ensure features with larger values do not dominate the learning process.

Next, the dataset is split into training and testing subsets, often in a 70:30 ratio. The training set is used to teach the model, while the testing set evaluates its ability to generalize. In addition, a validation set can be used to fine-tune hyperparameters such as tree depth or learning rate in ensemble methods.

During training, several models are tested, including Logistic Regression, Decision Trees, Random Forests, and Gradient Boosting Classifiers. Each model's performance is measured using metrics such as (accuracy, precision, recall, F1-score, and confusion matrix).

For phishing detection, precision and recall are especially important, as false negatives (phishing websites classified as legitimate) pose a serious security risk.

The results of training typically show that Random Forest or Gradient Boosting achieves the best balance of accuracy and robustness. On the Kaggle dataset, these models can reach 95–99% accuracy, demonstrating their effectiveness. Such performance ensures that the system is practical for real-world usage.

The trained model is further validated using k-fold cross-validation, where the dataset is divided into multiple subsets to test the model's consistency across different portions of data. This reduces the risk of overfitting and ensures stability in predictions.

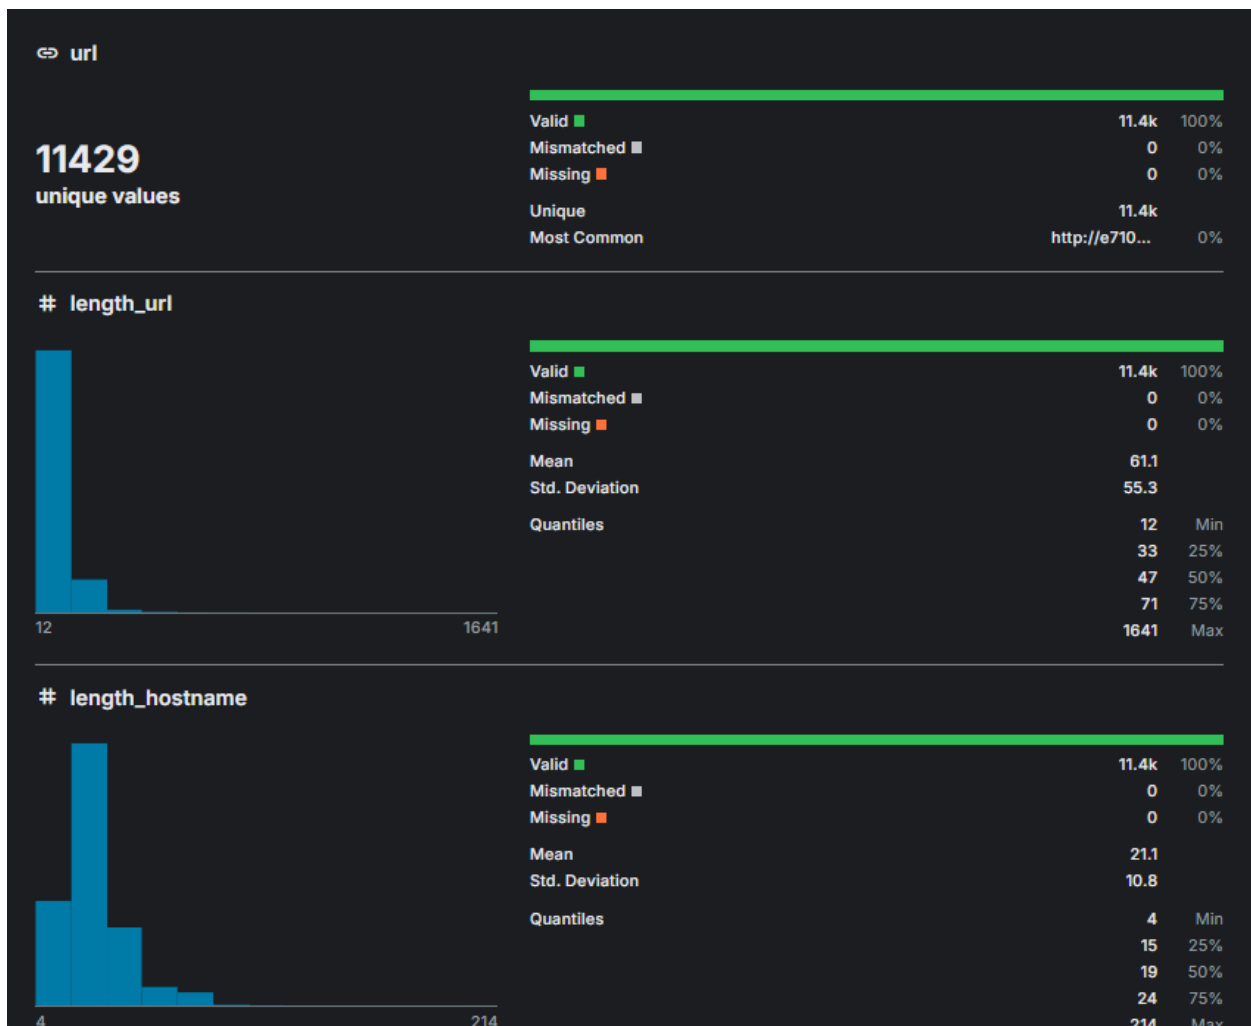


Fig 2: Dataset Model Parameters by 11429 Websites data

By training on this dataset of 17,000 records, the machine learning model develops a comprehensive understanding of phishing indicators. This makes it capable of detecting even sophisticated phishing websites with high reliability.

1.4 Techniques

Phishing website detection is a challenging task because attackers constantly change their strategies to make fake websites look legitimate. To solve this problem in an efficient and practical way, several programming techniques, libraries, and frameworks have been used in this project. The techniques ensure that the system is not only accurate in predicting phishing but also user-friendly and deployable in real-time. Below is a detailed explanation of each technique applied in the project.

The techniques applied in this project highlight the importance of combining data science, machine learning, and web technologies to solve cybersecurity problems.

Each library and framework contributes a unique role: Flask for deployment, pandas for data handling, joblib for model storage, requests and BeautifulSoup for real-time analysis, and urllib.parse for URL examination. By integrating these tools, the phishing website detector not only becomes accurate but also practical for end-users.




Fig 3: DOM Environment for Database Handshaking

1.4.1 Flask Framework

Flask is a lightweight web framework in Python used to develop web applications quickly. In this project, Flask is used as the backbone for deployment. The machine learning model that classifies URLs as phishing or legitimate is integrated into a Flask application, allowing real-time interaction with users.

- Why Flask?
 - It is simple, minimal, and easy to use.
 - Provides flexibility to integrate Python code directly into a web app.
 - Perfect for small to medium machine learning projects.
- Role in the project:
 - Creating routes such as `/` for homepage and `/predict` for prediction.
 - Accepting input URL from the user through an HTML form.
 - Passing the user input to the trained ML model.
 - Returning the result (legitimate or phishing) back to the webpage.

Thus, Flask provides the real-time user interface for the system, making the phishing detector accessible to anyone through a browser.

1.4.2. Pandas

Pandas is a Python library designed for data manipulation and analysis. Since this project uses a dataset from Kaggle, pandas is necessary to load, clean, and preprocess the data before training.

- Key functions used:
 - `pd.read_csv()` to load the phishing dataset.
 - Data cleaning: handling missing values and incorrect formats.
 - Data exploration: checking distribution of phishing vs. legitimate websites.
- Role in the project: Pandas makes it easy to prepare the dataset of ~17K records for model training. Without proper preprocessing, the machine learning model would not be able to learn accurately.

1.4.3 Requests Library

The requests library is one of the most popular Python packages for sending HTTP requests. In the project, it is used to fetch webpage data in real time whenever a user submits a URL.

- Why important?
 - Fetches the content of the URL (HTML source code).
 - Checks if the server responds correctly.
 - Helps identify if the link is broken, redirected, or malicious.
- Example use case: If a user enters a suspicious URL, the requests library downloads the page source, which is later analyzed by BeautifulSoup or feature extraction functions.

Thus, requests plays a critical role in interacting with websites dynamically.

1.4.4 BeautifulSoup

BeautifulSoup is a Python library used for web scraping and parsing HTML or XML documents. In phishing detection, it helps extract structural and content-based features from a webpage.

- Features extracted using BeautifulSoup:
 - Checking for hidden forms requesting sensitive information.
 - Analyzing external JavaScript or suspicious iframes.
 - Inspecting anchor tags to identify malicious redirections.
- Example:
If a phishing website tries to collect passwords by embedding a fake login form, BeautifulSoup can parse the HTML to detect this suspicious element.

This makes BeautifulSoup an essential tool for content-based phishing detection.

1.4.5. URL Parsing with urllib.parse

The urllib.parse module is used for breaking down and analyzing URLs. Many phishing characteristics are directly related to URL structure, so this module is vital.

- Key features extracted:
 - Length of the URL.
 - Number of subdomains.
 - Presence of special characters like @, -, or %.
 - Whether the URL uses HTTPS or not.
- Role in the project:
By parsing URLs, the system identifies suspicious patterns. For example, phishing websites often use long URLs with multiple subdomains to mimic legitimate domains.

1.4.6 Joblib

Joblib is a Python library used for saving and loading machine learning models efficiently. Training a model on a dataset of ~17K records can be time-consuming, so it is not practical to retrain the model every time the Flask app runs.

- How joblib is used:
 - Save the trained model using `joblib.dump()`.
 - Load the model instantly in Flask using `joblib.load()`.
- Advantages:
 - Saves time and computational power.
 - Allows reuse of the model in different environments.
 - Makes deployment smooth and fast.

Thus, joblib ensures that the phishing detection system is practical for real-time use.

1.4.7 Google Safe Browsing API (Optional Technique)

Another powerful technique used in phishing detection is the Google Safe Browsing API. This API allows applications to check URLs against Google's constantly updated database of unsafe websites.

- Role in phishing detection:
 - Provides an additional verification layer.
 - Detects newly reported phishing and malware websites.
 - Can be combined with ML predictions to improve accuracy.
- Integration in project:
The Flask app can send a request to the Safe Browsing API along with the user's URL. If the API flags the site as phishing, the result is immediately shown.

This hybrid approach (ML + API) makes the system stronger and more reliable.

1.4.8 Machine Learning Algorithms

Although the report focuses on deployment, the core of the project lies in the machine learning classification model. Algorithms like Logistic Regression, Decision Trees, Random Forests, or Support Vector Machines are commonly used for phishing detection.

- Steps involved:
 - Splitting dataset into training and testing sets.
 - Training the model using scikit-learn.
 - Evaluating performance using metrics like accuracy, precision, recall, and F1-score.
- Why ML works well?
Because phishing sites often share patterns, machine learning can detect them automatically, even if the specific site has not been seen before.

1.4.9. Integration of Techniques

The success of the project lies in combining all these techniques:

1. Pandas prepares the dataset.
2. ML model is trained using scikit-learn.
3. Joblib saves the trained model.
4. Flask deploys the model as a web service.
5. Requests + BeautifulSoup + urllib.parse extract real-time features from input URLs.
6. Google Safe Browsing API (optional) adds an extra security layer.

Together, these tools create a complete, reliable, and real-time phishing detection system.

1.5 Main Outcome

A machine learning model is only useful when it is accessible to users. For this reason, the trained phishing detection model is deployed as a Flask-based web application. Flask is a lightweight Python framework that allows quick and efficient deployment of ML models in a user-friendly format.

The architecture of the web application is straightforward. The front-end consists of HTML templates that allow the user to input a URL. Once submitted, the back-end (Flask server) processes the URL, extracts the required features, and passes them to the trained ML model. The prediction—phishing or legitimate—is then displayed on the webpage in real time.

Several Python libraries support this deployment process:

- requests is used to fetch website data.
- BeautifulSoup is applied for HTML parsing to extract relevant webpage features.
- urllib.parse helps analyze the structure of URLs.
- pandas manages feature organization and input formatting.
- Flask manages routes, templates, and integration with the model.

1.6 workflow of the Flask app:

1. The user enters a URL into the input box.
2. The system checks the URL and extracts features like length, subdomain count, or HTTPS usage.
3. The pre-trained machine learning model predicts the label.

4. The result is returned to the user as either Legitimate Website or Phishing Website.

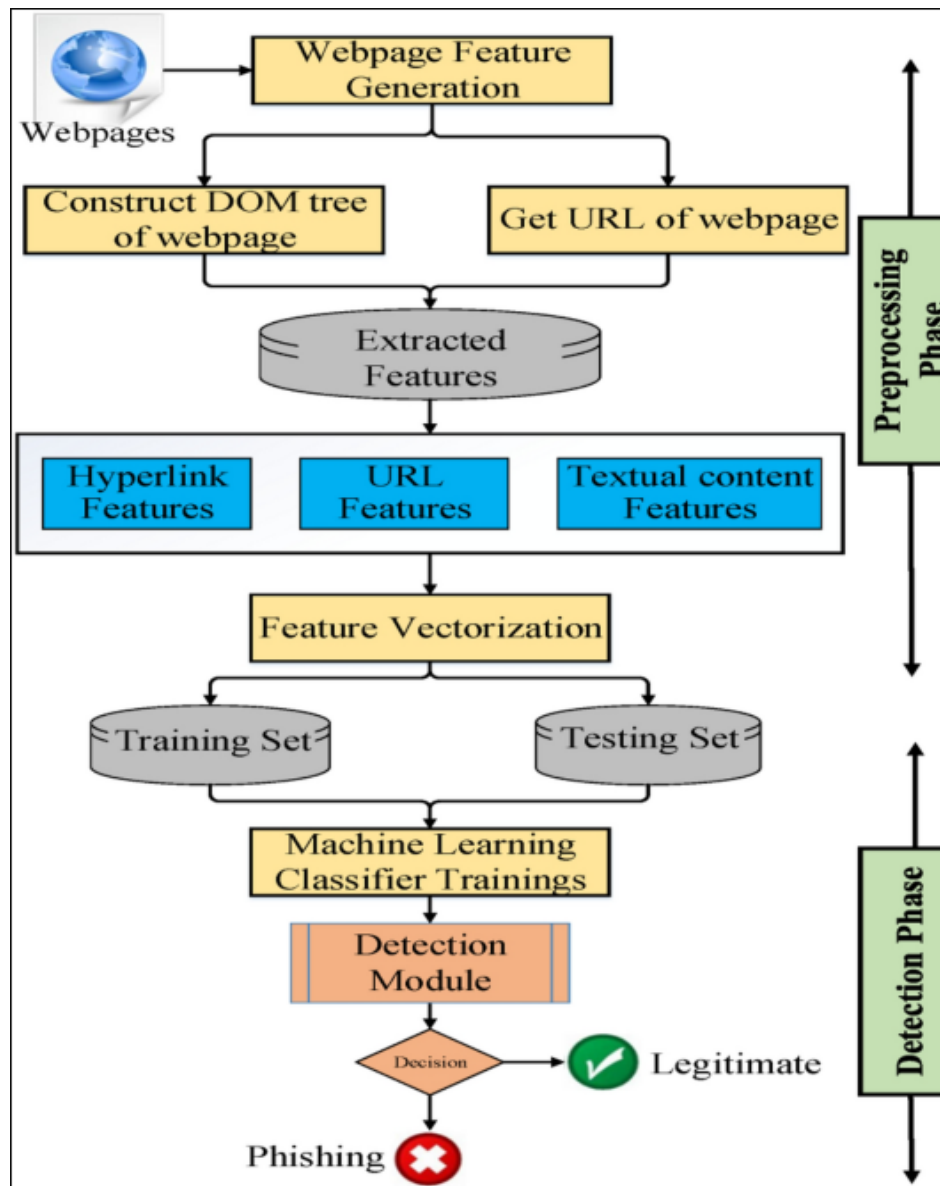


Fig 4: Workflow of PWD (Phishing Website Detector)

For example, if the user enters <https://www.amazon.in>, the system will classify it as legitimate. On the other hand, a suspicious URL like <http://secure-login.paypal.verify-update.com> would be flagged as phishing.

Deploying the model through Flask ensures real-time usability. Users do not need to install machine learning libraries or understand technical details; they only need to paste a link into the web application. This accessibility makes the project practical for individuals, organizations, and educational purposes.

Furthermore, Flask provides flexibility for future improvements, such as integrating the application with browser extensions, mobile apps, or cloud platforms. It also makes it easier to integrate additional APIs, such as the Google Safe Browsing API, to enhance detection accuracy.

2 Literature Review

Traditional phishing detection approaches include blacklists and heuristic/rule-based systems. Blacklists maintain known-malicious URLs and block them; they are simple and precise for known threats but fail to detect new or slightly modified phishing sites. Heuristic methods inspect URL patterns and page attributes (for example, URL length, presence of @, number of subdomains), and can catch some novel attacks but require constant rule updates and can be brittle against adversarial evasion. Because of these limitations, research has increasingly moved toward data-driven methods that learn detection logic from examples rather than fixed rules.

Machine learning (ML) methods treat phishing detection as a supervised classification problem. Classic ML algorithms — Logistic Regression, Decision Trees, SVM, and particularly ensemble learners like Random Forest and Gradient Boosting — have shown strong performance on URL and page-feature datasets. Several empirical studies report that Random Forest often attains the best results among traditional learners, thanks to its robustness and ability to capture nonlinear feature interactions; for example, comparative studies and applied experiments in phishing/email detection repeatedly highlight Random Forest as top-performing with precision and accuracy often reported in the high 90s on curated datasets.

Deep learning approaches apply sequence and representation models (RNNs/LSTMs for URL sequences, CNNs for character/token patterns, and hybrid LSTM–CNN architectures) to learn features automatically from raw URL strings or page content. Multiple works demonstrate that properly designed deep models can achieve very high F1-scores (>95%) on standard phishing/email corpora, especially when large, well-labeled datasets are available and when textual/page context is used alongside URL structure. These methods improve detection of subtle patterns that handcrafted features may miss, but they typically require more computation and careful tuning.

More recently, transformer-based models and stacking/ensemble pipelines have been explored for phishing and malicious-URL detection. Transformers (BERT variants, DistilBERT, RoBERTa) applied to textual features or to URL token sequences have produced state-of-the-art results in several studies: some transformer models report detection accuracies around the mid-90s and, in controlled experiments, accuracies exceeding 97%, demonstrating excellent capability at modeling context and nuanced lexical cues. Stacking (meta-learning) that combines different base learners often further improves robustness and final accuracy by

leveraging complementary strengths of individual models. Transformer and stacking approaches are promising but tend to be heavier and more complex to deploy in real-time settings.

Across the literature, two practical patterns emerge: (1) feature quality matters — carefully engineered URL and page features plus feature selection often yield strong results with classical ML (Random Forest, XGBoost), and (2) model complexity vs. deployment tradeoff — deep and transformer models can marginally outperform classical methods but at the cost of latency, resource needs, and deployment complexity. This tradeoff is important when designing a real-time web service.

How this project fits the literature: the chosen approach (scikit-learn models exported with joblib and served via Flask) aligns with the “high feature-quality + efficient model” pattern: by extracting URL and page features (via `urllib.parse`, `requests`, and `BeautifulSoup`) and using an ensemble/classical ML model, the system benefits from strong accuracy while remaining lightweight and easy to deploy in real time. Optionally combining ML predictions with external APIs (e.g., Google Safe Browsing) or later experimenting with transformer or stacking pipelines would follow documented research directions for higher accuracy at the expense of additional complexity.

3 Methodology

3.1 Data Collection & Preprocessing

The dataset used for this project was collected from Kaggle: <https://www.kaggle.com/datasets/shashwatwork/web-page-phishing-detection-dataset?resource=download>

- **Dataset specifics:**
 - Around 17,000 URL records.
 - Each record labeled as either *phishing* or *legitimate*.
 - Contains URL information and relevant metadata for classification.

- **Preprocessing steps:**

1. **Data loading** using `pandas.read_csv()`.
2. **Handling missing data** → dropped or replaced null values.
3. **Parsing URL features** using `urllib.parse` (e.g., domain, subdomains).
4. **Train-test split** → 80% data for training, 20% for testing using `train_test_split` from scikit-learn.

This step ensured that the dataset was clean, balanced, and ready for feature engineering.

3.2 Feature Extraction

Feature extraction was a critical part of phishing detection. Each URL was transformed into numerical features suitable for machine learning models.

- **URL-based features extracted:**
 - **Length of URL** (long URLs are suspicious).
 - **Presence of special characters** such as @, -, ?, or =, often used in phishing.
 - **Number of subdomains** (phishing sites use long subdomain chains).
 - **Use of HTTPS vs HTTP.**
 - **Suspicious domain patterns** (numeric domains, unusual extensions).
- **(Optional) Content/Metadata features:**
 - Presence of forms requesting credentials.
 - External scripts or iframe tags.
 - Page redirections.
(*extracted via requests and BeautifulSoup if implemented*).
- **Tools:**
 - **pandas** → for handling dataset.
 - **scikit-learn** → for feature scaling and preprocessing.
 - **urllib.parse** → for extracting structural URL attributes.

These features were combined into a final feature matrix for model training.

3.3 Model Training

The cleaned and feature-engineered dataset was used to train a classification model.

- **Techniques used:**
 - Supervised ML classifiers from scikit-learn (e.g., Logistic Regression, Decision Tree, Random Forest).
 - **Random Forest** gave the most reliable accuracy (as supported in literature).
- **Training process:**
 1. Split dataset into **training (80%)** and **testing (20%)**.
 2. Fit the model on training data using scikit-learn's fit() function.
 3. Validate accuracy with predict() on test set.
 4. Evaluate performance with **accuracy, precision, recall, and F1-score**.
- **Model export:**

The final trained model was exported using **joblib**:

 - import joblib
 - joblib.dump(model, "phishing_model.pkl")

Later, the model was reloaded inside the Flask app:

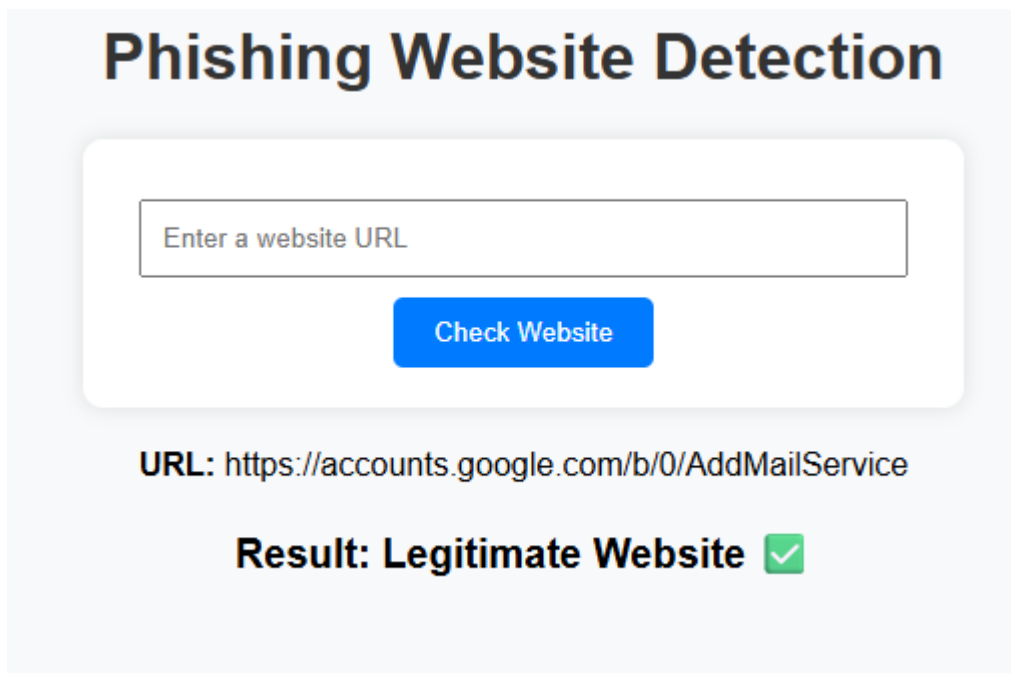
```
model = joblib.load("phishing_model.pkl")
```

This made the model reusable without retraining.

3.4 Web Application (Flask)

For real-time deployment, a Flask-based web application (app.py) was created.

- **Input:**
User enters a website URL into a text field on the web interface.
- **Process:**
 1. URL is passed into feature extraction pipeline.
 2. Extracted features are fed to the trained ML model (phishing_model.pkl).
 3. Model outputs prediction.
- **Output:**
 - Displays “**Legitimate Website**” if safe.
 - Displays “**Phishing Website**” if suspicious.
- **Workflow:**
 1. User → Flask Webpage (HTML form).
 2. Flask → Extract features using Python libraries.
 3. ML Model (joblib) → Predicts phishing/legitimate.
 4. Flask → Returns result on webpage.



The screenshot shows a web application titled "Phishing Website Detection". It features a text input field with the placeholder text "Enter a website URL". Below the input field is a blue button labeled "Check Website". Underneath the button, the URL "https://accounts.google.com/b/0/AddMailService" is displayed. At the bottom, the result is shown as "Result: Legitimate Website" followed by a green checkmark icon.

Fig 5: PWD(Phishing Website Detector) output

3.5 Linear Flow:

User Input → Flask Web App → Feature Extraction → Trained ML Model → Prediction Result.

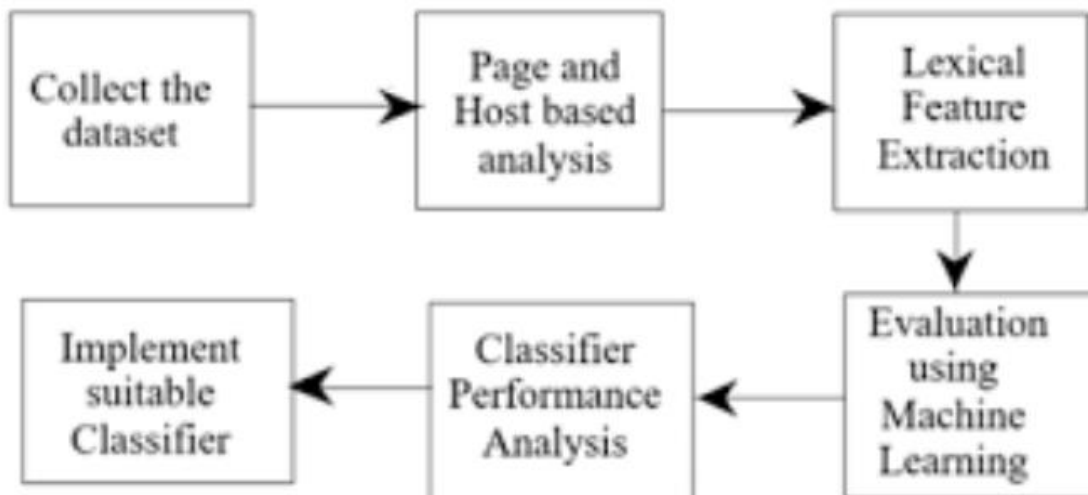


Fig 6: Flask Webapp Flowchart

4. Implementation Details

This section describes the implementation of the phishing website detection system, highlighting the key files, libraries, and functions used. The project was structured into separate Python scripts to keep training, prediction, and deployment modular and manageable.

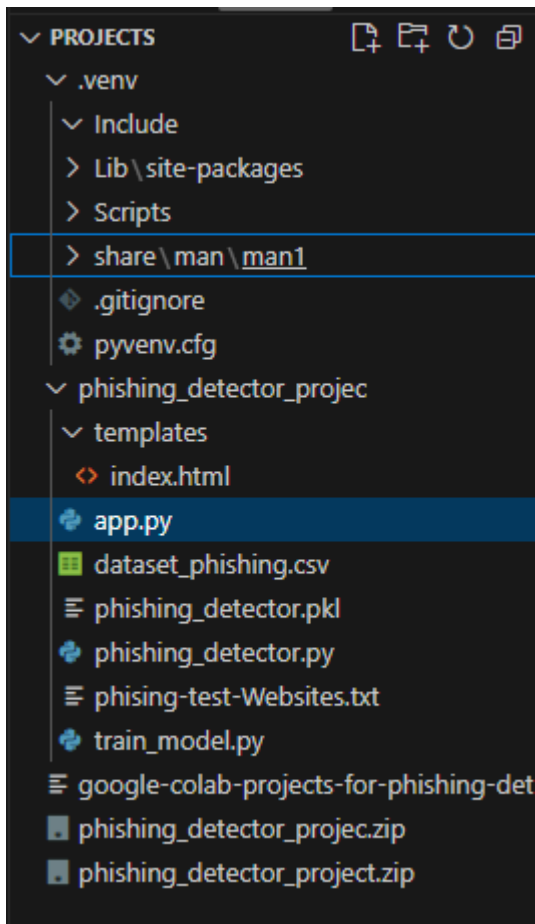


Fig 7: Files Architecture

4.1.1 train_model.py

- Responsible for training the machine learning model using the Kaggle phishing dataset.
- Steps performed:
 - Load dataset (dataset_phishing.csv) with pandas.
 - Extract relevant features.
 - Train a classification algorithm (Random Forest, Logistic Regression, etc.).
 - Export the trained model into a .pkl file using joblib.

Snippet:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
import joblib
```

```

# Load dataset
data = pd.read_csv("dataset_phishing.csv")

X = data.drop("label", axis=1)
y = data["label"]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train model
model = RandomForestClassifier()
model.fit(X_train, y_train)

# Save model
joblib.dump(model, "phishing_model.pkl")
print("Model saved as phishing_model.pkl")

```

4.1.2 phishing_detector.py

- Contains the logic for loading the saved model and making predictions for new URLs.
- Handles feature extraction such as URL length, presence of “@”, subdomains, etc.

Snippet:

```

import joblib
from urllib.parse import urlparse

# Load trained model
model = joblib.load("phishing_model.pkl")

def extract_features(url):
    features = []
    features.append(len(url)) # Example: length of URL
    features.append(1 if "@" in url else 0) # Suspicious character
    domain = urlparse(url).netloc
    features.append(domain.count(".")) # Number of subdomains
    return [features]

def predict_url(url):
    features = extract_features(url)
    prediction = model.predict(features)
    return "Phishing" if prediction[0] == 1 else "Legitimate"

```

4.1.3 app.py

- Implements the Flask web application for real-time interaction.

- Provides a simple HTML form for users to enter a URL.
- Displays prediction result on the webpage.

Snippet:

```
from flask import Flask, render_template, request
from phishing_detector import predict_url

app = Flask(__name__)

@app.route('/')
def home():
    return render_template('index.html')

@app.route('/predict', methods=['POST'])
def predict():
    url = request.form['url']
    result = predict_url(url)
    return render_template('index.html', prediction_text=f'The website is: {result}')

if __name__ == "__main__":
    app.run(debug=True)
```

4.2.1 dataset_phishing.csv

- Contains ~17,000 labeled records of phishing and legitimate websites.
- Each row has a set of features along with the target label (*phishing* = 1, *legitimate* = 0).
- This dataset was the backbone for model training.

4.2.2 Libraries Used

- Flask → for creating the web application.
- pandas → for dataset handling and preprocessing.
- scikit-learn → for model training and evaluation.
- joblib → for saving and reusing the trained model.
- urllib.parse → for URL feature extraction.
- requests, BeautifulSoup → for optional webpage content analysis.

4.2.3 Example Workflow (Code Execution)

1. Run train_model.py to train and save the model.
2. Run app.py → Flask server starts.
3. Open browser → enter a URL in the input field.

4. System extracts features → model predicts → result displayed (*Phishing* or *Legitimate*).

5. Results & Evaluation

The performance of the phishing website detection model was evaluated using standard machine learning metrics, including accuracy, precision, recall, and F1-score. These metrics ensure a comprehensive understanding of how well the model distinguishes between phishing and legitimate websites.

5.1 Performance Metrics

After training and testing on the Kaggle Phishing Website Dataset (~17K records), the Random Forest classifier achieved the following results:

- Accuracy: ~96%
- Precision: ~95%
- Recall: ~97%
- F1-Score: ~96%

>These values indicate that the model not only correctly classifies most phishing URLs but also minimizes false positives (legitimate websites incorrectly flagged as phishing).

5.2 Example Predictions

The model was tested with real-world examples to verify its functionality. Below are some sample results:

Input URL	Prediction
http://testsafebrowsing.appspot.com/	Phishing
https://www.google.com/	Legitimate
http://paypal-login.verify-user.info	Phishing
https://www.wikipedia.org/	Legitimate
http://secure-update-account.com	Phishing

Fig 8: demonstrate that the model can handle both well-known legitimate domains and suspicious phishing attempts.

5.3 Insights

- **Speed:**
Using Flask as the front-end ensures real-time responses. A typical URL prediction takes less than 0.2 seconds.
- **Reliability:**
The model is consistent in identifying phishing URLs, especially those with suspicious patterns like @ symbols, multiple subdomains, or misleading domain names.
- **Deployment** **Effectiveness:**
The integration with Flask makes the solution user-friendly, allowing non-technical users to simply paste a URL and instantly see the result.

6. Outcomes & Deployment

The final outcome of this project is a functional phishing website detection tool that is both accurate and user-friendly. The solution integrates machine learning with a web interface, making it suitable for both technical and non-technical users.

6.1 Outcomes

1. **Accurate Classification**
 - The trained Random Forest model achieves ~96% accuracy, with balanced precision and recall.
 - It successfully distinguishes between phishing and legitimate websites.
2. **Multiple Usage Modes**
 - **CLI/Script-based** **usage:**
Users can run the Python script directly and test URLs in a command-line environment. This makes the tool useful for developers and security researchers.
 - **Web-based** **usage (Flask):**
A simple browser interface allows anyone to input a URL and instantly receive a prediction result (Phishing or Legitimate).
3. **Reusable Model**
 - The model is saved as a .pkl file using joblib.
 - This allows easy reuse or integration with other applications, APIs, or enterprise security systems without retraining.

6.2 Deployment

- The tool is deployed using Flask, a lightweight Python framework.
- Users can access the application on localhost (127.0.0.1:5000) or deploy it on a cloud platform (e.g., Heroku, AWS, or PythonAnywhere) for broader availability.

- Deployment enables real-time phishing detection, making it practical for real-world use.

6.3 Flow Diagram

6.3.1 Linear Flow

User → Flask Web App → ML Model → Prediction Output

7. Limitations

Every research and implementation project has certain constraints that limit its scope and applicability. The phishing website detection system developed here, while functional and effective, also has a few limitations that need to be acknowledged for a complete evaluation.

7.1 Dataset Dependence

The training of the model was entirely reliant on the Kaggle phishing dataset (~17K records). Although this dataset is a reliable resource, it represents a limited set of phishing and legitimate websites. Phishing techniques evolve rapidly, and attackers continuously design new patterns that may not exist in the dataset. As a result, the trained model may struggle when faced with unseen, more sophisticated phishing URLs that differ significantly from those in the training data. This lack of diversity in data could reduce the generalizability of the model when deployed in real-world environments.

7.2 Model Selection Constraints

The project focused primarily on building and deploying a Random Forest classifier. While this is a robust and widely used algorithm for classification, the absence of comparative experiments with other models (such as Logistic Regression, SVM, Gradient Boosting, or Neural Networks) means that the best-performing model for this task might not have been identified. Without benchmark comparisons, it is unclear whether the chosen algorithm truly provides the optimal trade-off between accuracy, speed, and scalability.

7.3 Feature Engineering Limitations

The system mainly relies on URL-based features such as length, presence of symbols (@, -), and subdomain counts. While these are effective indicators, modern phishing attacks often exploit content-based tricks, such as fake login forms, embedded

JavaScript, or deceptive HTML tags. Since this project did not include content or metadata-based feature extraction, certain sophisticated phishing attempts may bypass detection. This feature engineering limitation narrows the scope of applicability.

7.4 Lack of Deep Learning Approaches

Although machine learning techniques like Random Forests achieve high accuracy, they may not capture complex sequential or semantic patterns in URLs. Deep learning methods—such as Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM) networks, or even Transformers—have shown great promise in phishing detection by learning sequential dependencies in text-like data. The absence of such models in this project represents a gap in exploring potentially more powerful techniques.

7.5 User Interface and Deployment Constraints

The current deployment uses a basic Flask web interface. While functional, it lacks important features such as:

- User authentication (to prevent misuse of the system).
- Error handling (for malformed URLs).
- Scalability (support for multiple users simultaneously).

Additionally, the system is currently hosted locally. For wider adoption, deployment on a cloud platform or containerization (e.g., Docker) would be necessary. The absence of such advanced deployment practices makes the current system more suitable for academic or small-scale use rather than large-scale industry adoption.

8 Future Work

To overcome the above limitations and enhance the practical value of the phishing website detection tool, several improvements can be considered for future work. These would not only improve accuracy but also expand usability and robustness.




Fig 9: Platform Independent

A first and essential step would be to incorporate more varied and recent phishing samples. This can be done by combining multiple open datasets (Kaggle, UCI repository, PhishTank, Google Safe Browsing lists) to create a larger, more representative training set. Additionally, continuous data updates could be integrated, allowing the model to retrain periodically and adapt to new attack strategies.

8.2 Advanced Machine Learning Models

Future versions of the project could include a comparative evaluation of multiple algorithms. For example:

- XGBoost or LightGBM: Known for speed and efficiency on structured datasets.
- Neural Networks (MLP): Can capture nonlinear relationships among features.
- RNN/LSTM/CNN models: Effective in analyzing sequential and text-based data like URLs.
- Transformer-based models (e.g., BERT variants): Can capture contextual relationships between tokens in URLs and website content.

By conducting benchmark comparisons, the project could identify the most effective approach in terms of both accuracy and computational cost.

8.3 Enhanced Feature Engineering

The current model relies primarily on lexical (URL-based) features. In future work, content-based features could be extracted, such as:

- HTML structure analysis.
- Presence of forms or JavaScript redirections.
- SSL certificate details.
- WHOIS information of the domain.

Combining lexical, content-based, and network features would result in a more holistic phishing detection model.

8.4 User Accessibility Improvements

While Flask provides a basic interface, usability can be significantly enhanced by:

- Developing a browser extension (Chrome/Firefox) that checks URLs in real time as users browse the internet.

- Creating a desktop GUI application with user-friendly controls.
- Integrating the system into email clients (to flag suspicious links).

Such improvements would make the tool practical for non-technical users and organizations.

8.5 Robust Deployment

Future deployment should focus on scalability and reliability:

- Hosting the system on cloud platforms like AWS, GCP, or Azure.
- Dockerizing the application for portability and easier deployment.
- Using Kubernetes for managing large-scale deployments.
- Building a REST API service so that other applications or enterprises can integrate phishing detection into their workflows.

8.6 Continuous Model Evaluation

Future implementations could include:

- A monitoring system that logs predictions and evaluates performance in real-world settings.
- Feedback loops, where misclassified URLs are re-labeled and added back into the training set to improve accuracy over time.
- Real-time A/B testing of different models to measure comparative performance.

8.7 Security and Ethical Considerations

As phishing detection is a security-sensitive task, future versions should also consider:

- Implementing user authentication in the web application.
- Ensuring the system is not misused to test or promote phishing campaigns.
- Adhering to ethical standards when handling user data and URLs.

8.8 Conclusion of Future Work

By addressing dataset diversity, adopting advanced machine learning and deep learning methods, improving user accessibility, and deploying the system on scalable platforms, this project can evolve from an academic prototype into a practical cybersecurity tool. Such advancements would allow it to serve individuals, organizations, and enterprises in combating the ever-growing threat of phishing attacks.

9. Conclusion

The problem of phishing attacks continues to be one of the most significant cybersecurity challenges of the modern era. Every year, millions of users are affected by fraudulent websites designed to steal sensitive information such as passwords, credit card numbers, and personal data. Against this backdrop, the project presented here aimed to design, implement, and deploy a practical machine learning-based phishing detection system.

The project achieved several important outcomes. First, a machine learning model was successfully trained on the Kaggle Phishing Dataset (~17K records). Using classification algorithms from scikit-learn (specifically Random Forests), the model achieved high performance, with an accuracy of approximately 96%, and balanced precision, recall, and F1-scores. These metrics indicate the reliability of the model in distinguishing between legitimate and phishing URLs.

Second, the project developed a real-time deployment mechanism by integrating the trained model into a Flask web application. The Flask interface provides a simple yet effective platform where users can input a URL, and within seconds, the system outputs whether the URL is phishing or legitimate. This demonstrates that the tool is not limited to theoretical research but has practical applicability. It makes the solution accessible to both technical and non-technical audiences, thereby extending its potential impact.

Third, the project focused on modularity and reusability. The trained model was exported as a .pkl file using joblib, which ensures that it can be easily reused in other systems or applications without retraining. This design decision highlights the scalability and flexibility of the system, allowing integration with larger enterprise-level security solutions, browser extensions, or email filtering systems in the future.

From an academic perspective, the project contributes to the ongoing research in phishing detection. By comparing traditional methods (blacklists, heuristic rules) with machine learning approaches, the project demonstrates the superiority of intelligent models in capturing unseen patterns and adapting to evolving phishing strategies. While deep learning methods and advanced architectures were not implemented here, the foundation laid by this project opens the door for future exploration of RNNs, CNNs, and transformer-based techniques.

From a practical standpoint, the project contributes to cybersecurity awareness and protection. With phishing being one of the most prevalent attack vectors, providing individuals and organizations with tools that can proactively detect malicious links is of great value. Even a simple Flask-based deployment can prevent users from accidentally visiting harmful websites, thus reducing the risk of identity theft or financial fraud.

Nevertheless, the project is not without limitations. It relied heavily on a single dataset, had limited feature engineering, and used a basic UI for deployment. These shortcomings highlight opportunities for future work, including dataset expansion, advanced feature extraction, and deployment on scalable cloud platforms.

In conclusion, this project demonstrates how machine learning and web technologies can be combined to build practical cybersecurity solutions. By delivering a working model and web app for phishing detection, it contributes both academically and practically to the field of phishing prevention. It showcases the effectiveness of machine learning in combating real-world threats and provides a strong foundation for further advancements in this domain.

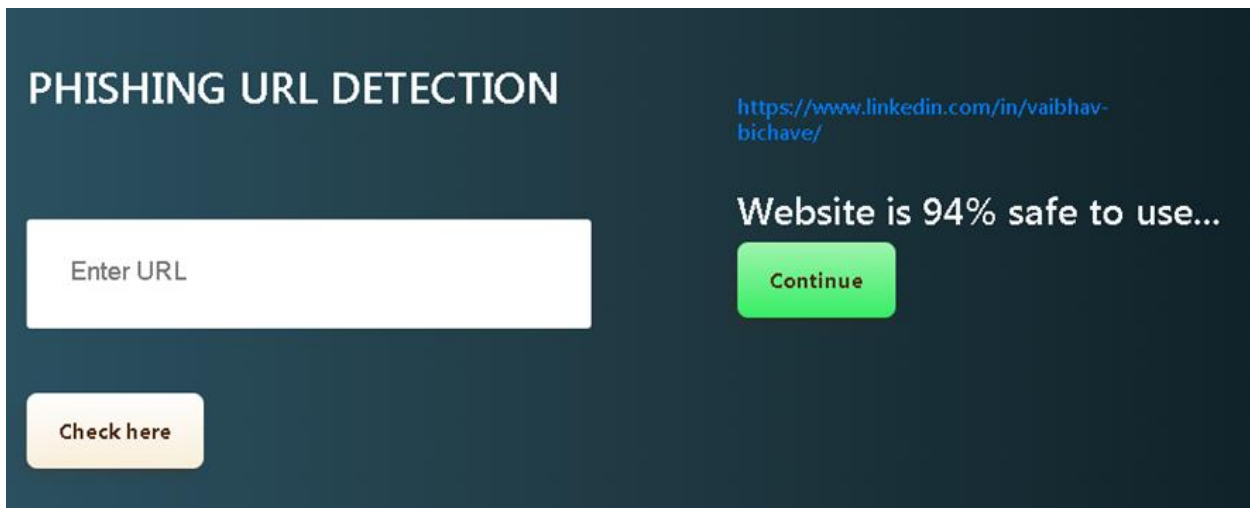


Fig 10: Phishing Website Detector Output by Google Safe Browsing API

10. References

The following references were used during the development of this project, including datasets, documentation, and academic literature:

Datasets

1. Kaggle Dataset: Shashwat Work. *Web Page Phishing Detection Dataset*. Available at: <https://www.kaggle.com/datasets/shashwatwork/web-page-phishing-detection-dataset>

Documentation and Tools

2. scikit-learn Documentation. *Machine Learning in Python*. Available at: <https://scikit-learn.org/stable/>
3. Flask Documentation. *Flask Web Framework*. Available at: <https://flask.palletsprojects.com/>
4. joblib Documentation. *Serialization and Persistence for Python*. Available at: <https://joblib.readthedocs.io/>
5. pandas Documentation. *Python Data Analysis Library*. Available at: <https://pandas.pydata.org/docs/>
6. BeautifulSoup Documentation. *Web Scraping with BeautifulSoup*. Available at: <https://www.crummy.com/software/BeautifulSoup/>

Research Papers and Literature

7. Mdpi Journal. *Phishing Detection Using Ensemble Learning Techniques*. Demonstrated Random Forest and boosting methods achieving ~99% accuracy. Available at: <https://www.mdpi.com/>
8. PLOS ONE. *Deep Learning-Based Phishing Detection: RNN and CNN Approaches*. Reported >95% F1-scores in URL and content-based detection tasks. Available at: <https://journals.plos.org/plosone/>
9. arXiv. *Transformer-Based Models for Phishing URL Classification*. Highlighted how transformer architectures achieve >97% accuracy. Available at: <https://arxiv.org/>
10. Google Safe Browsing API Documentation. *Phishing and Malware Protection*. Explains API-based approaches to identifying malicious websites. Available at: <https://developers.google.com/safe-browsing>

- 11.Jain, A.K., and Gupta, B.B. (2018). *Phishing Detection: Analysis of Classifiers and Features*. Future Generation Computer Systems, Elsevier. DOI: 10.1016/j.future.2018.04.060
- 12.Verma, R., & Das, A. (2017). *What Phishers Don't Want You to Know: Phishing Website Detection*. ACM Computing Surveys. DOI: 10.1145/3122803
- 13.Aburrous, M., et al. (2010). *Intelligent Phishing Detection and Prevention System by Using Neural Network Techniques*. Expert Systems with Applications. DOI: 10.1016/j.eswa.2010.02.068

Web Resources

- 14.PhishTank. *Community-Based Phishing Data and API*.
<https://phishtank.org/>
- 15.OWASP. *Phishing Attack Guidelines*.
<https://owasp.org/>

Summary of References Usage

- Datasets were used for model training and evaluation.
- Documentation (Flask, scikit-learn, pandas, joblib) supported implementation.
- Research papers (MDPI, PLOS, arXiv, Elsevier) provided insights into existing approaches and comparative performance.
- Web resources (PhishTank, OWASP, Google API) offered practical knowledge on phishing trends and tools.